

## Background

The main code is inspired by :

<https://github.com/udacity/deep-reinforcement-learning/tree/master/ddpg-pendulum>

## Algorithms

In this project, a modified DDPG is in place to train the agent. The main idea of DDPG is shown below. Basically, there are two actor networks (local and target) to learn the policy, using (state, action). There are another two critic networks (local and target) to learn the Q-value function, using the action generated from actor network. Then use a soft-update to update target model for both actor and critic.

---

**Algorithm 1** Deep Deterministic Policy Gradient

---

```
1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi$ , empty replay buffer  $\mathcal{D}$ 
2: Set target parameters equal to main parameters  $\theta_{\text{targ}} \leftarrow \theta$ ,  $\phi_{\text{targ}} \leftarrow \phi$ 
3: repeat
4:   Observe state  $s$  and select action  $a = \text{clip}(\mu_{\theta}(s) + \epsilon, a_{\text{Low}}, a_{\text{High}})$ , where  $\epsilon \sim \mathcal{N}$ 
5:   Execute  $a$  in the environment
6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$ 
8:   If  $s'$  is terminal, reset environment state.
9:   if it's time to update then
10:    for however many updates do
11:      Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$ 
12:      Compute targets
```

$$y(r, s', d) = r + \gamma(1 - d)Q_{\phi_{\text{targ}}}(s', \mu_{\theta_{\text{targ}}}(s'))$$

```
13:   Update Q-function by one step of gradient descent using
```

$$\nabla_{\phi} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi}(s, a) - y(r, s', d))^2$$

```
14:   Update policy by one step of gradient ascent using
```

$$\nabla_{\theta} \frac{1}{|B|} \sum_{s \in B} Q_{\phi}(s, \mu_{\theta}(s))$$

```
15:   Update target networks with
```

$$\begin{aligned}\phi_{\text{targ}} &\leftarrow \rho \phi_{\text{targ}} + (1 - \rho) \phi \\ \theta_{\text{targ}} &\leftarrow \rho \theta_{\text{targ}} + (1 - \rho) \theta\end{aligned}$$

```
16:   end for
17: end if
18: until convergence

---


```


## Lessons learnt while tuning the hyper-parameters

1. Use normalization for neural network input (this is also suggested from original paper). This will help converge faster
2. Use sigmoid activation for critic output layer
3. Ornstein-Uhlenbeck process noise generation, using a smaller sigma helps to converge (use 0.05 here).
4. Using update interval from “benchmark implementation”: update the networks 10 times after every 20 timesteps.
5. Use clipping policy for critic model to avoid cliff jumping.
6. Add a noise decay for generating the action.

## (Hyper)parameters

Actor model	Neural network
	First linear layer: input 33 neurons, output 400 neurons. Activation function: Relu
	Second linear layer: input 400 neurons, output 300 neurons. Activation function: Relu
	Third linear layer: input 300 neurons, output 4 neurons. Activation function: tanh
	Note. A normalization is done before Relu from first to second layer.
Critic Model	Neural network
	First linear layer: input 33 neurons, output 400 neurons. Activation function: Relu
	Second linear layer: input 404 neurons, output 300 neurons. Activation function: Relu
	Third linear layer: input 300 neurons, output 1 neurons. Activation function: sigmoid
	Note. A normalization is done before Relu from first to second layer.
Replay buffer size	1000000.0
Batch size	128
Discount factor, Gamma	0.99
Actor model learning rate	0.001
Critic model learning rate	0.001
Weight decay	0
Soft update TAU	0.001
Learning internal	20
Learning num	10
Noise epsilon	1
Noise epsilon decay	0.999

With current strategy, after 300 episode, the score is steady to 31.9(see the plot below).



Next step

1. The neural network architecture could be fine-tuned.
2. Tune other hyper-parameters to achieve the goal faster.
3. Try out other methods, eg. REINFORCE, etc.